

Mizan.ai

RAG Pipeline Flowcharts

Two end-to-end flowcharts for Feature A's knowledge-base pipeline:

1. Offline / Ingestion Pipeline — runs locally, no Modal calls. Fetches source documents, extracts, chunks, handles tables and cross-references, embeds with a local BGE-M3, and writes to Qdrant + SQLite.
2. Online / Inference Pipeline — runs on Modal, on every user query. Language detection, jurisdiction check, retrieval with same-language-first + fallback, table/reference resolution, context capping, and answer generation with citation enforcement.

Companion to the RAG Data Source Strategy and Embeddings/Chunking/Pipeline Architecture documents.

Mizan.ai — Offline / Ingestion Pipeline (runs locally, no Modal calls)



